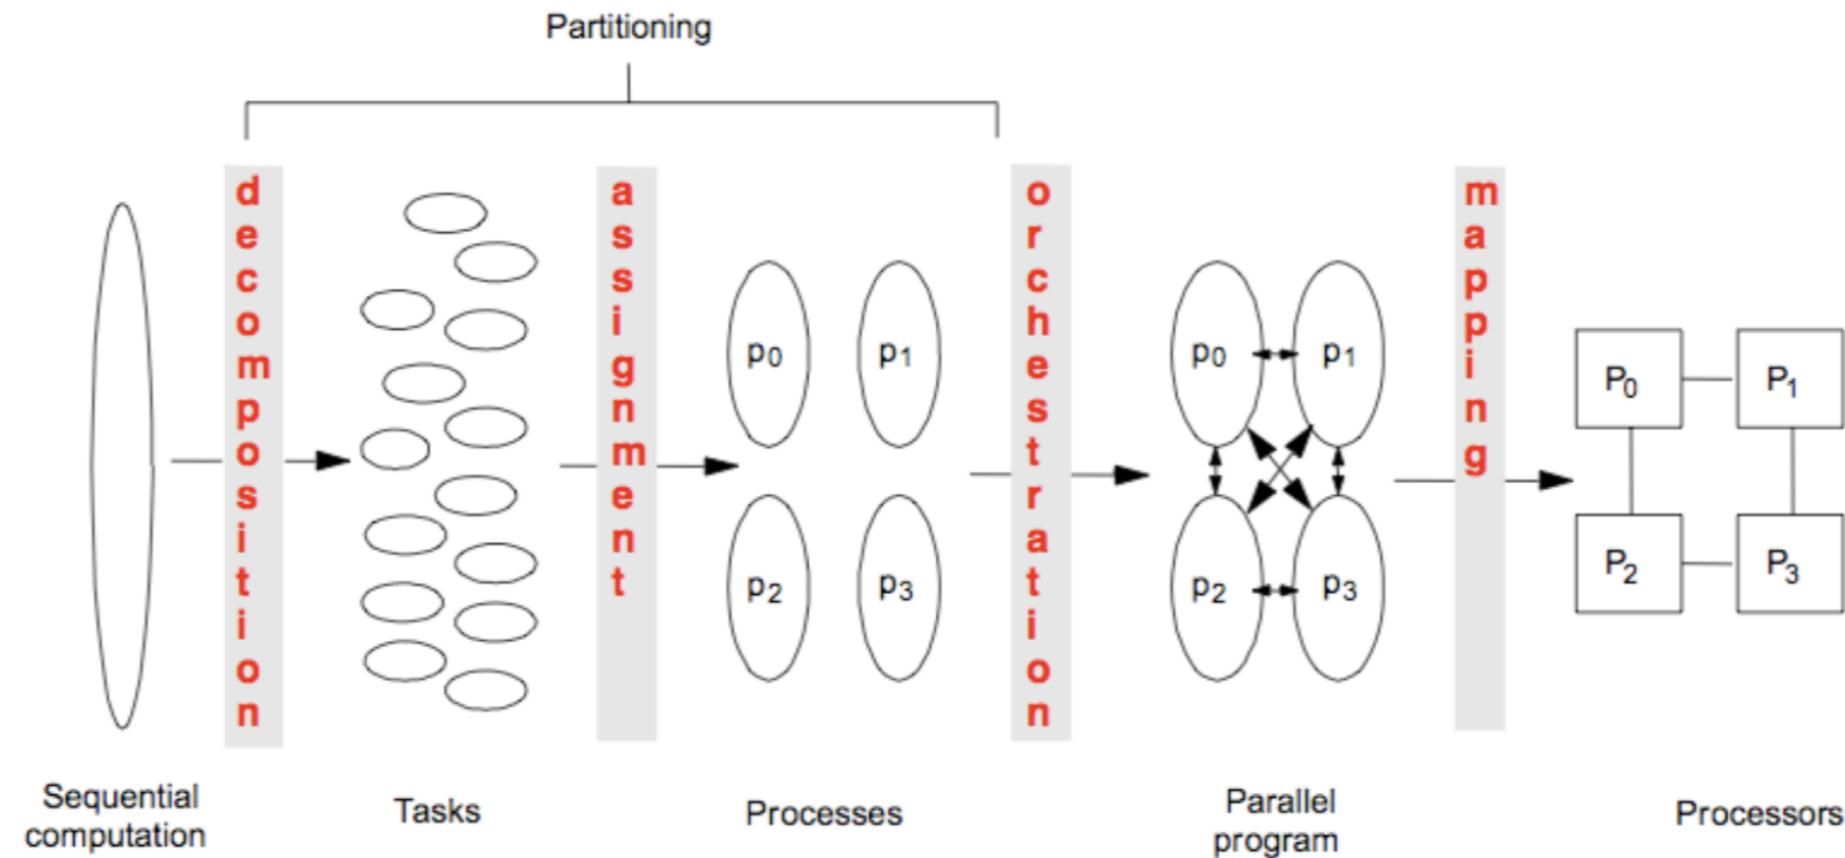


# Parallel Algorithm Design

# Parallel Algorithm Design

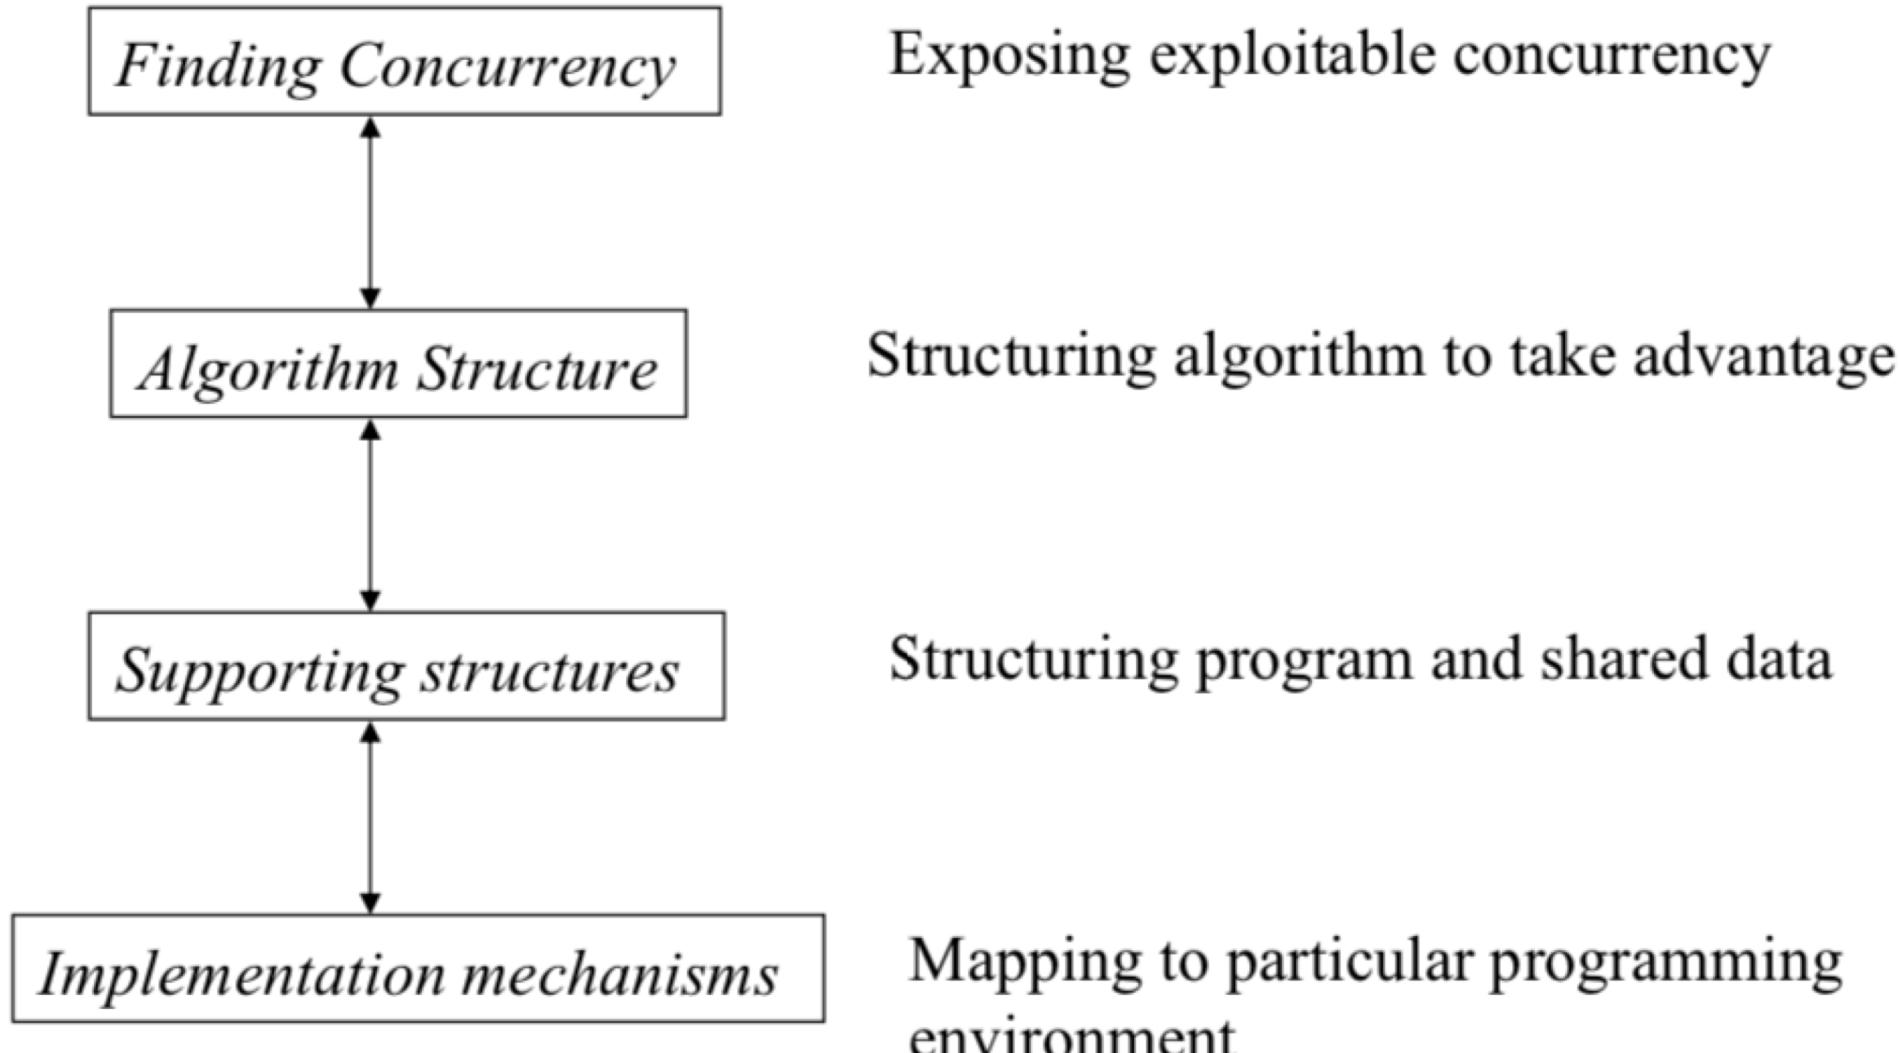


**1. Study problem or code**

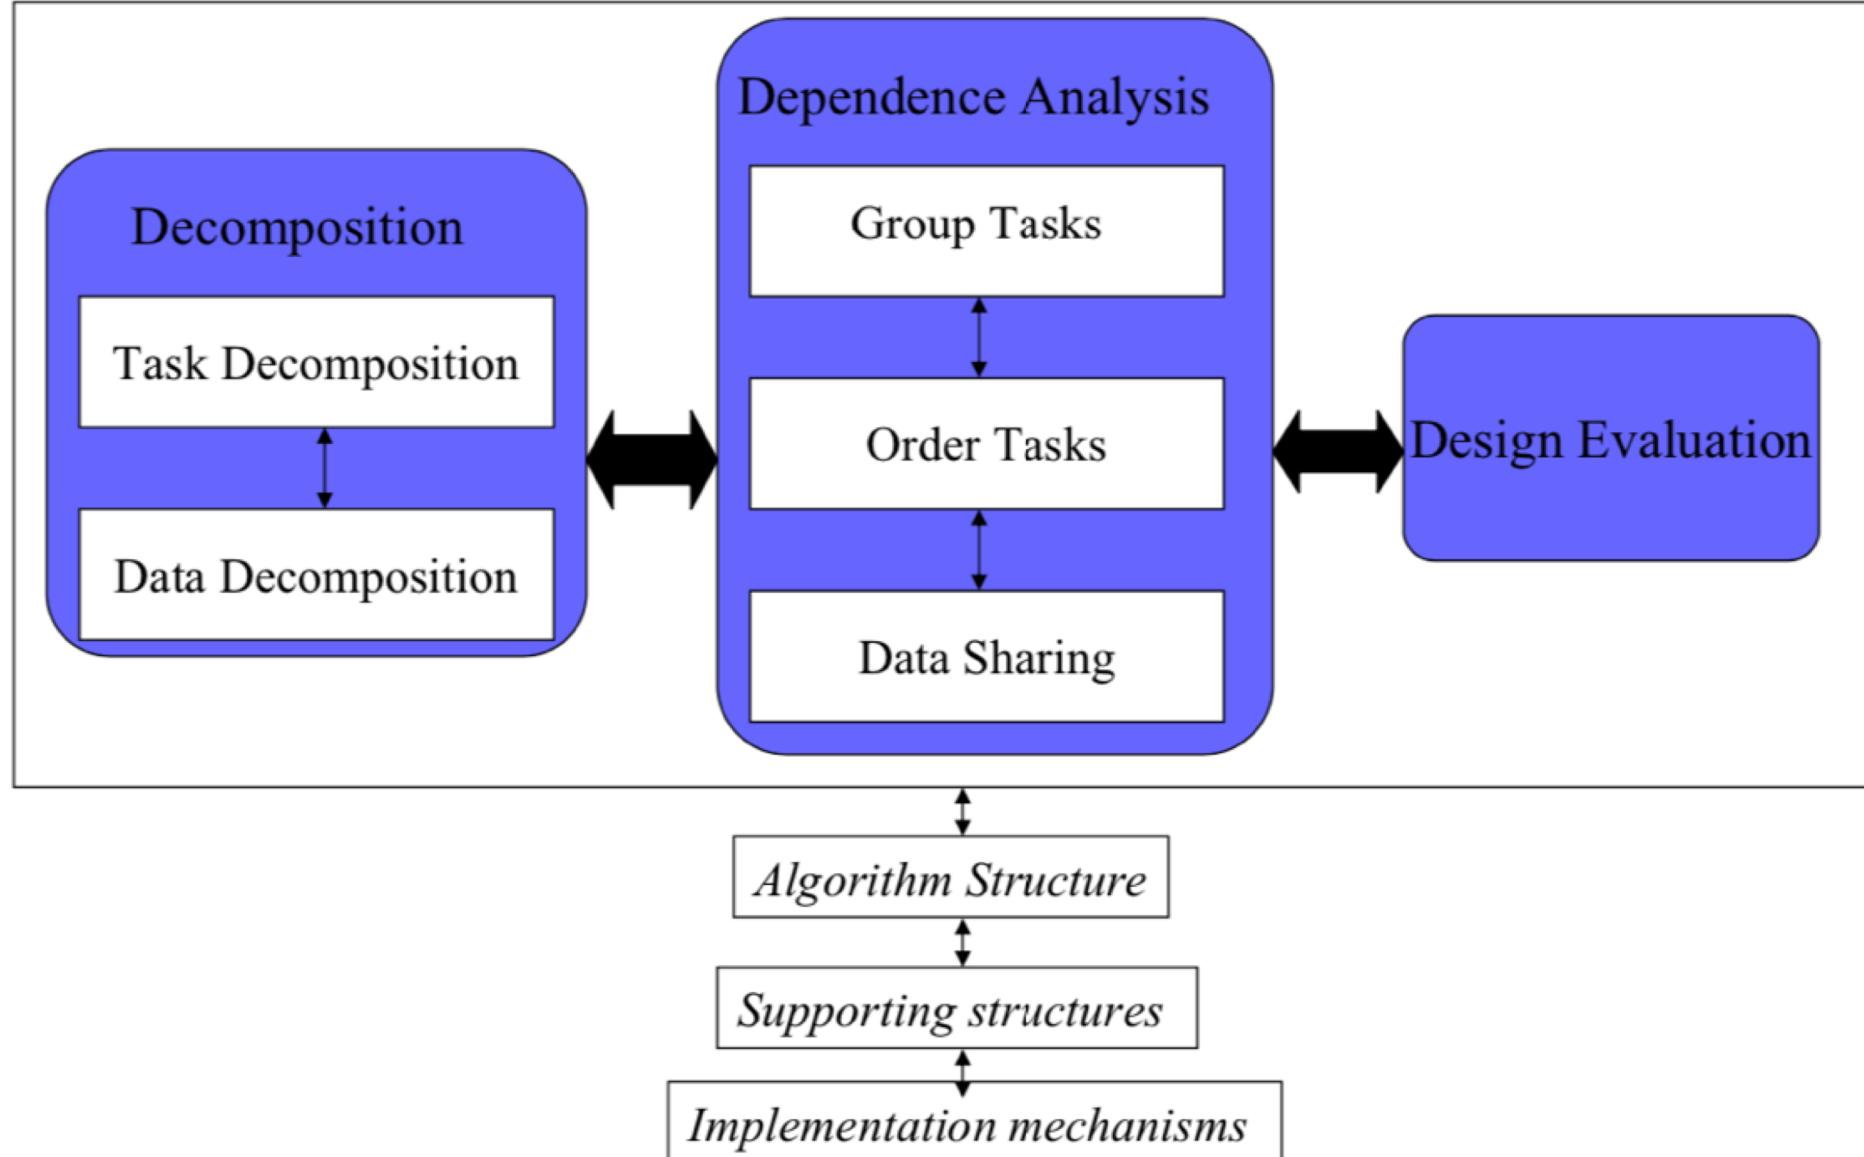
**2. Look for parallelism opportunities**

**3. Try to keep all cores busy doing useful work**

# Parallel Algorithm Design



# Finding Concurrency



# Decomposition Patterns

- Task decomposition: view problem as a stream of instructions that can be broken into sequences called tasks that can execute in parallel.
  - Key: Independent operations
- Data decomposition: view problem from data perspective and focus on how the can be broken into distinct chunks
  - Key: Data chunks that can be operated upon independently
- Task and data decomposition imply each other. They are different facets of the same fundamental decomposition

# Example

- Matrix multiplication
  - Task decomposition
    - Considering the computation of each element in the product matrix as a separate task
    - Performs poorly => group tasks pattern
  - Data decomposition
    - Decompose the product matrix into chunks, e.g., one row a chunk, or a small submatrix (or block) per chunk

# Dependency Analysis Pattern

- Group tasks: group tasks that have the same dependency constraints; identify which tasks must execute concurrently
  - Reduced synchronization overhead – all tasks in the group can use a barrier to wait for a common dependence
  - All tasks in the group efficiently share data loaded into a common on-chip, shared storage (Shard Memory)
  - Grouping and merging dependent tasks into one task reduces need for synchronization
- Order task pattern: identifying order constraints among task groups.
  - Control dependency: Find the task group that creates it
  - Data dependency: temporal order for producer and consumer relationship

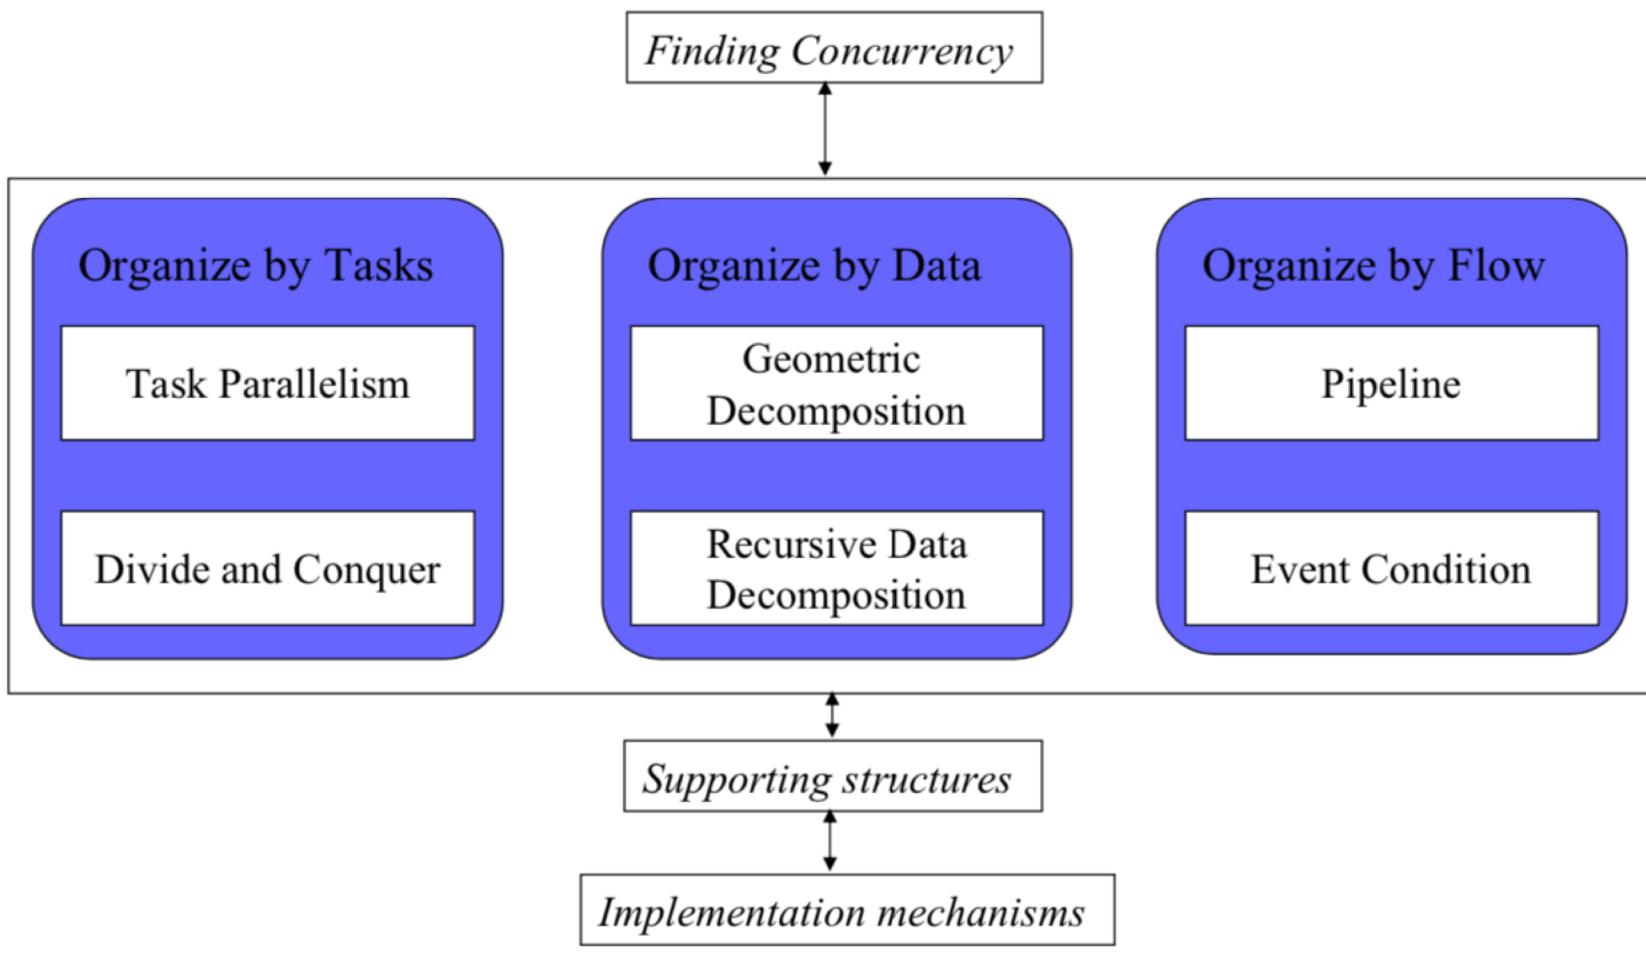
# Dependency Analysis Pattern

- Data sharing pattern: how data is shared among the tasks?
  - Read only: make own local copies
  - Effectively local: the shared data is partitioned into subsets, each of which is accessed (for read or write) by only one task a time.
  - Read-write: the data is accessed by more than one task.  
Need exclusive access mechanisms.
  - Example: the use of the shared memory among threads in a thread block.

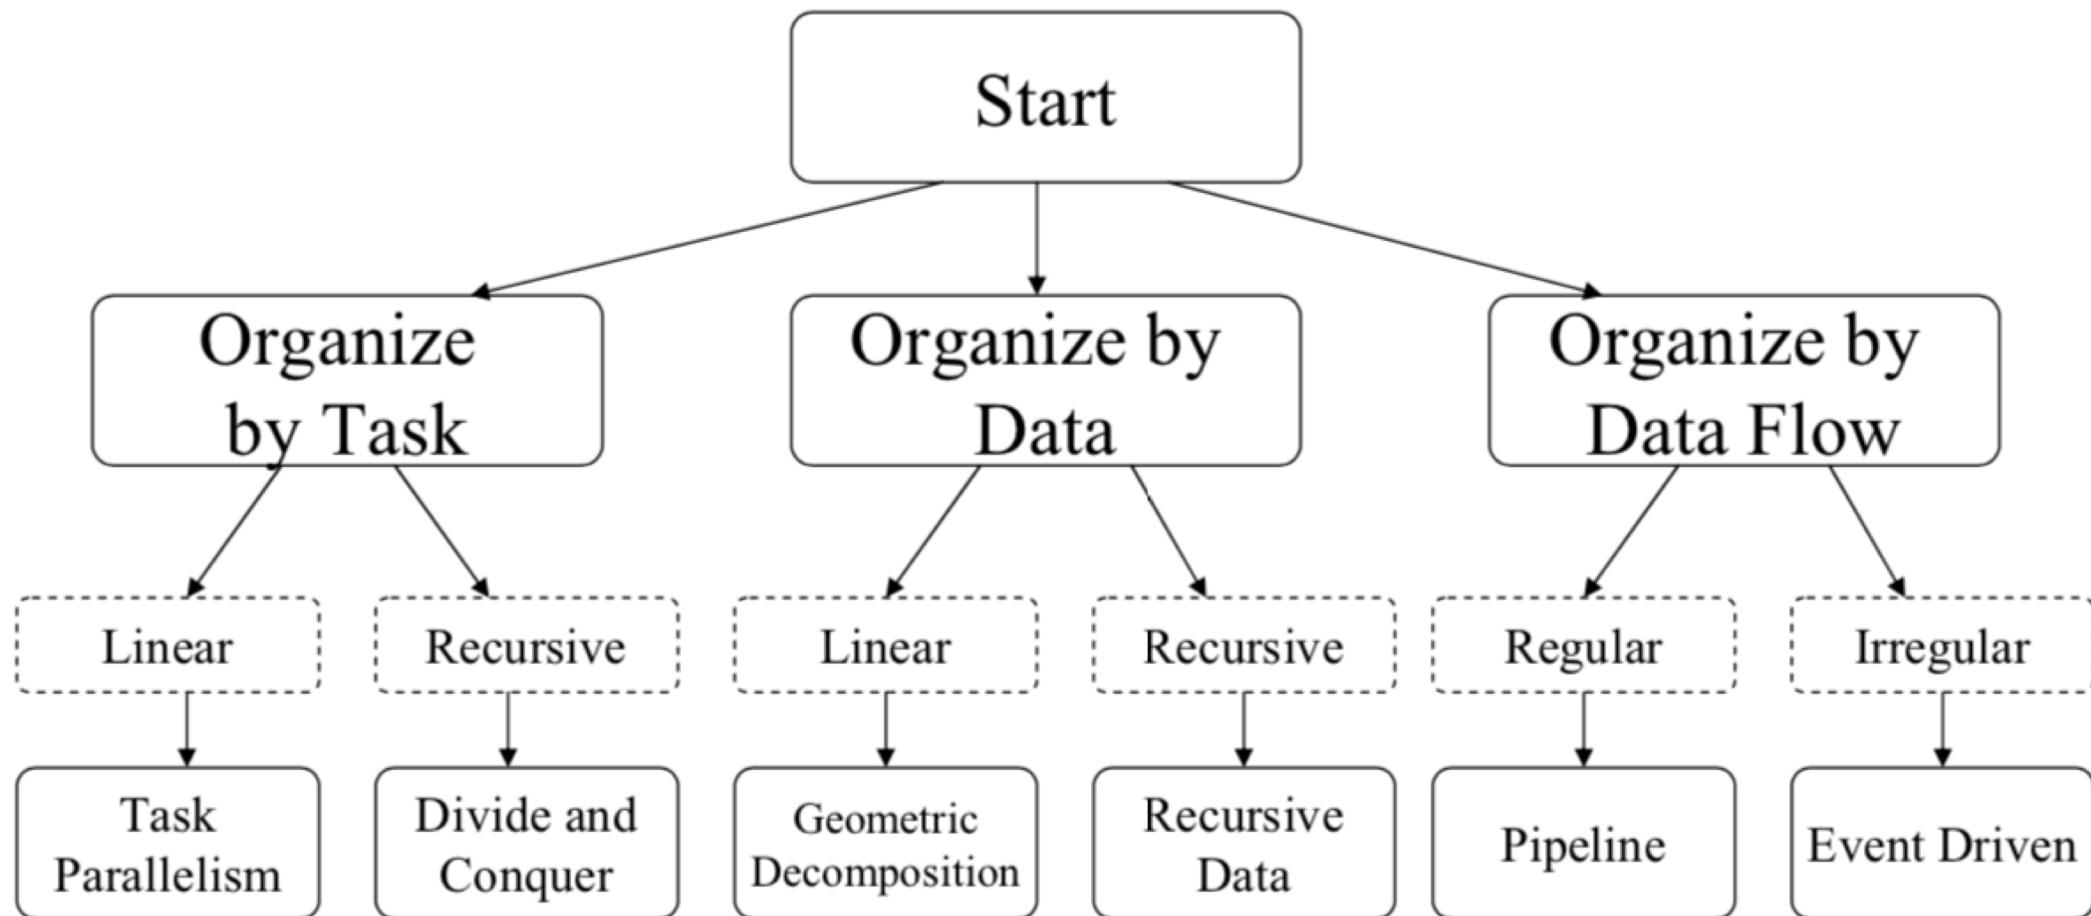
# Design Evaluation Pattern

- Whether the partition fits the target hardware platform?
- Key questions to ask
  - How many threads can be supported?
  - How many threads are needed?
  - How are the data structures shared?
  - Is there enough work in each thread between synchronizations to make parallel execution worthwhile?

# Algorithm Structure

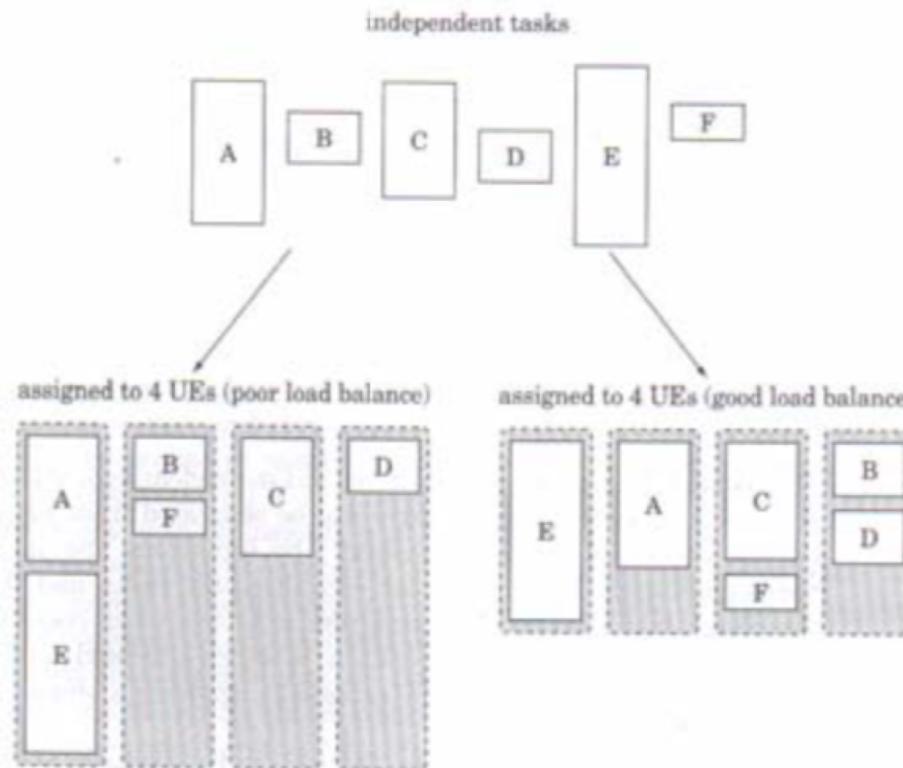


# Organizing Principle

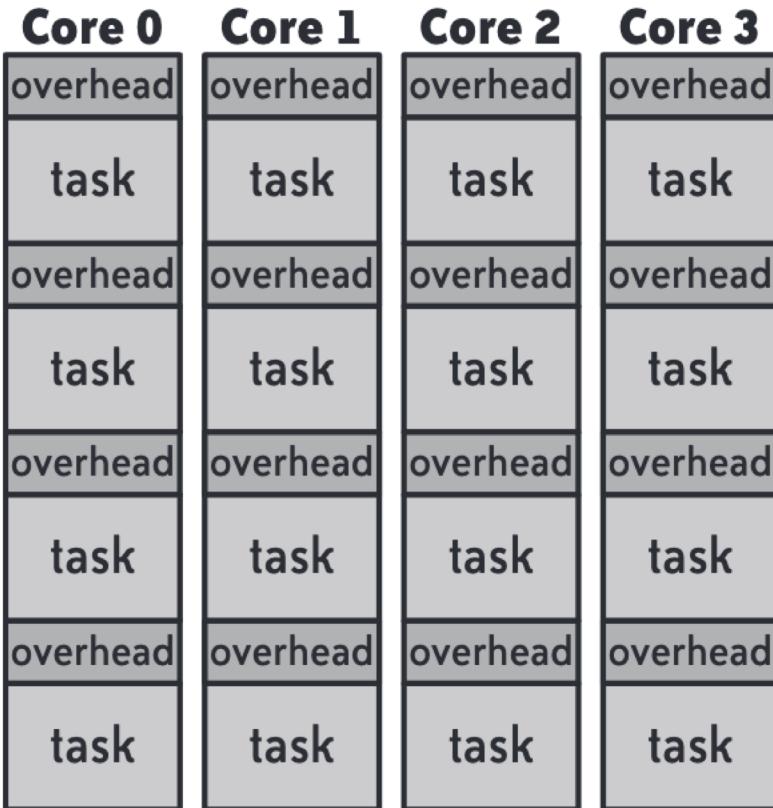


# Task Parallelism Pattern

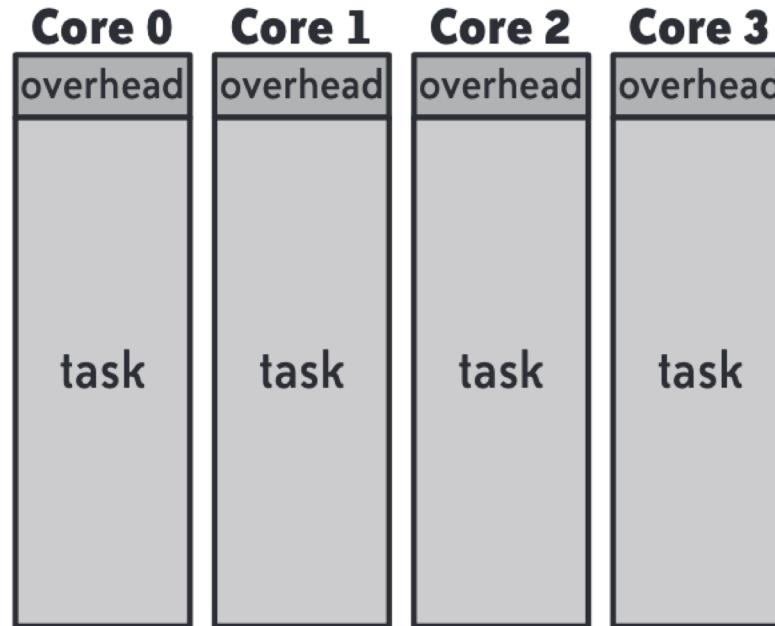
- After the problem is decomposed into a collection of tasks that can execute concurrently, how to exploit this concurrency efficiently?
- Load balancing



# Task Granularity Example



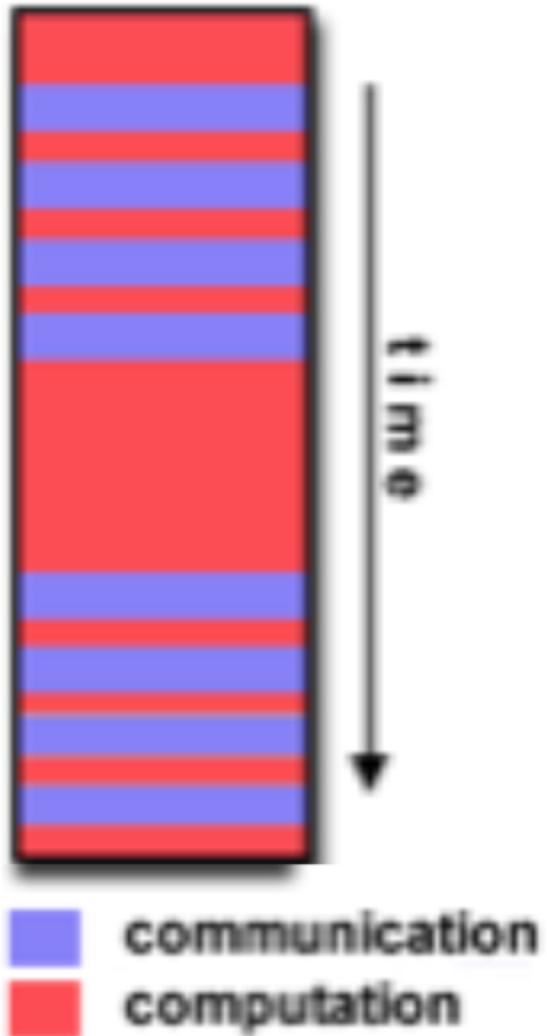
**(a) Fine-grained decomposition**



**(b) Coarse-grained decomposition**

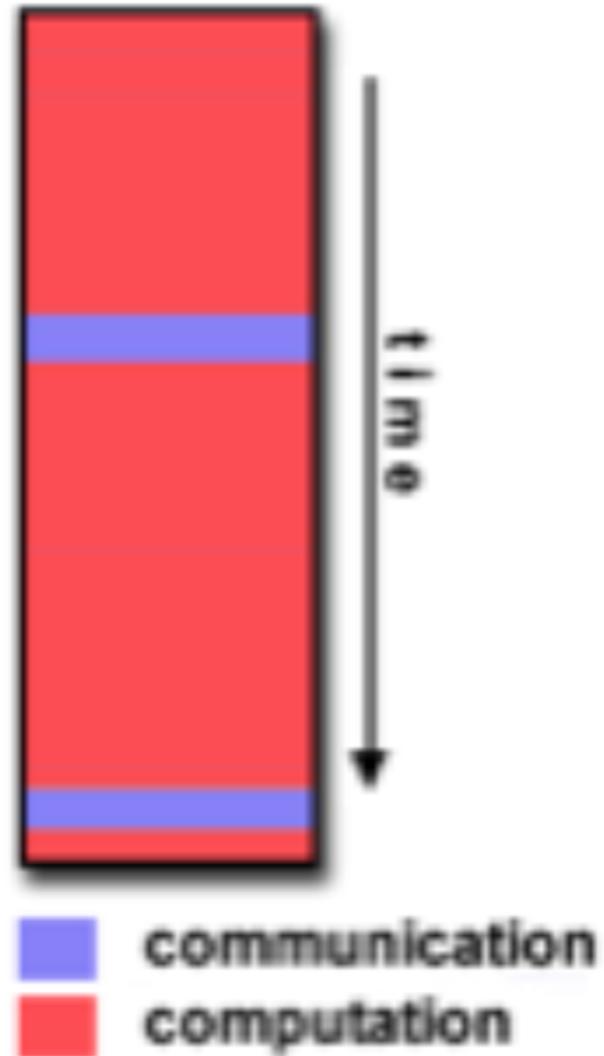
# Task Granularity Example

- Tasks execute little comp. between comm.
- Easy to load balance
- If ***too fine***, comm. may take longer than comp.



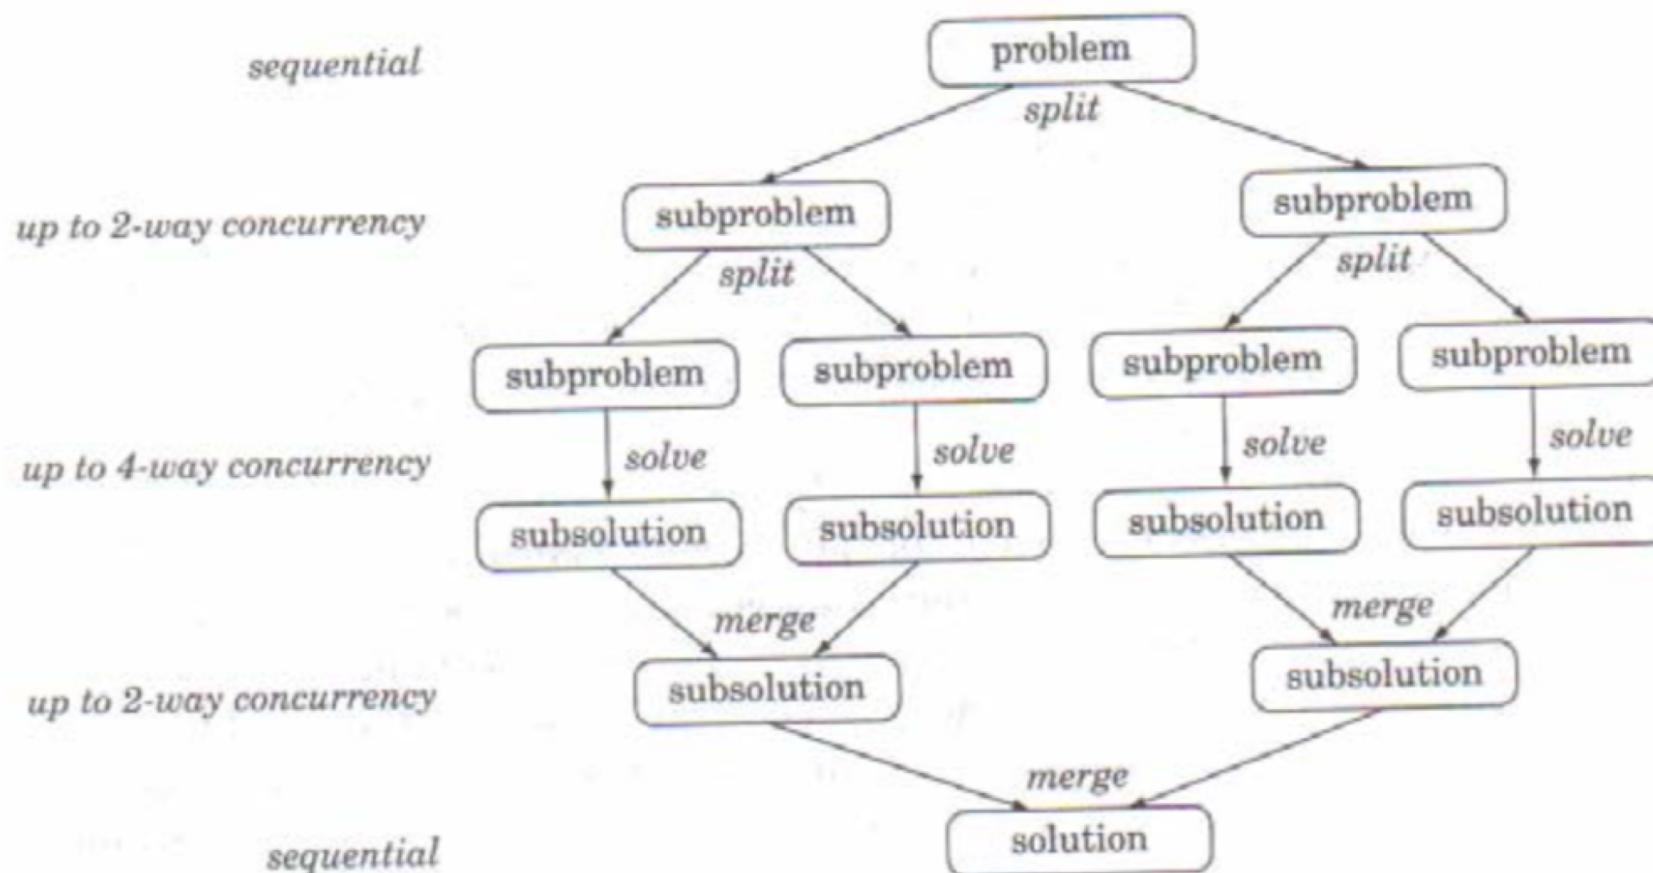
# Task Granularity Example

- Long computations between communication
- More opportunity for performance increase
- Harder to load balance



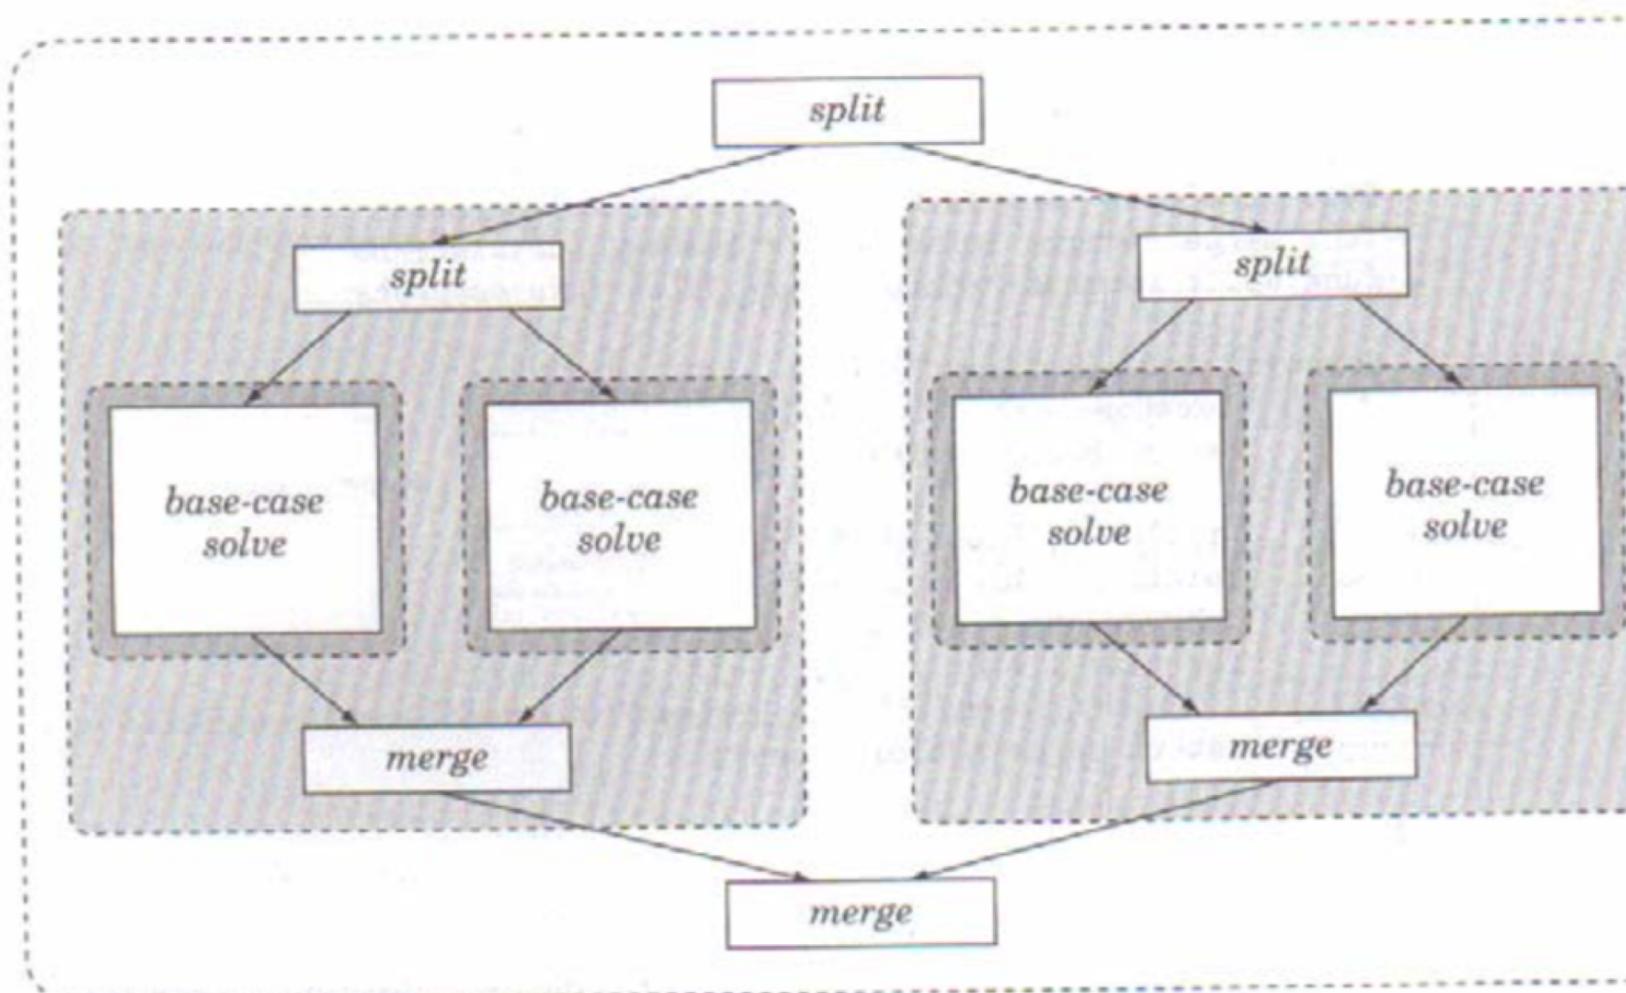
# Divide and Conquer Pattern

- If the problem is formulated using the sequential divide-and-conquer strategy, how to exploit the potential concurrency?



# Divide and Conquer Pattern

- Parallelization Strategy



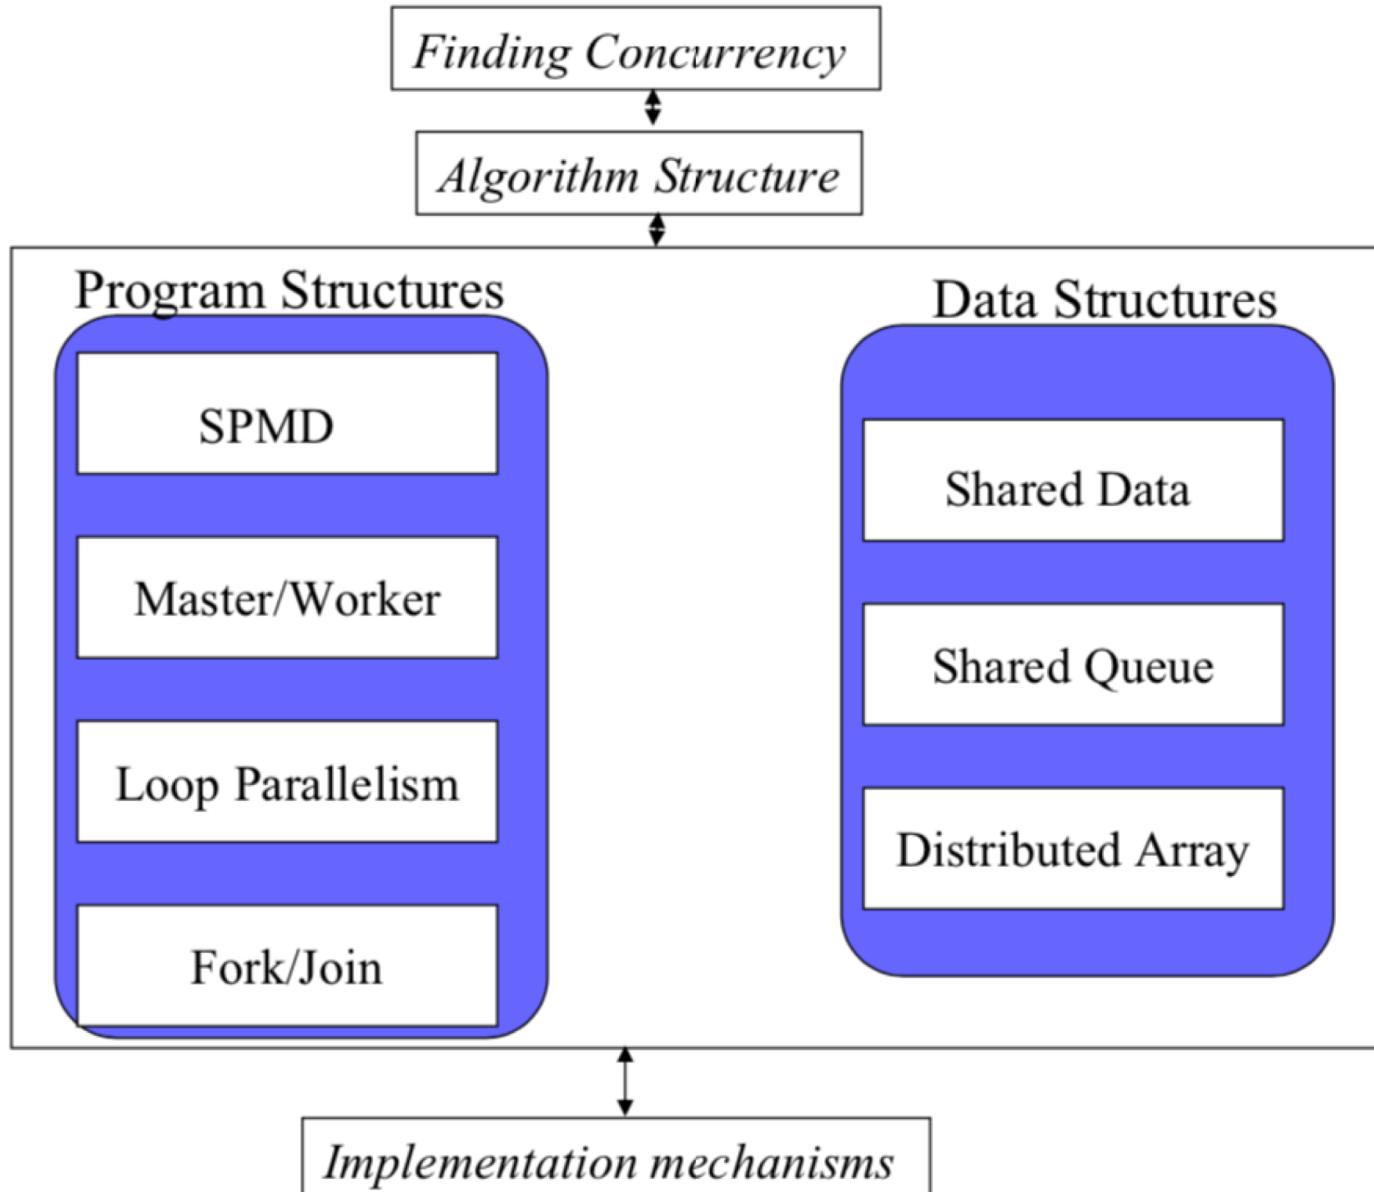
# Geometric Decomposition Pattern

- How to organize the algorithm after the data has been decomposed into concurrently updatable chunks?
- Decomposition to minimize the data communication and dependency among tasks
- Care needs to be taken when update non-local data, e.g., exchange operations

# Recursive Data Pattern

- Suppose the problem involves an operation on a recursive data structure that appears to require sequential processing. How to make the operations on these data structures parallel?
- Check whether divide-and-conquer pattern works
- If not, may need to transform the original algorithm.

# Supporting Structures



# Relationship between supporting program structure pattern and Algorithm structure pattern

	Task Parallel.	Divide/Conquer	Geometric Decomp.	Recursive Data	Pipeline	Event-based
SPMD	😊😊😊😊	😊😊😊	😊😊😊😊	😊😊	😊😊😊	😊😊
Loop Parallel	😊😊😊😊	😊😊	😊😊😊			
Master /Worker	😊😊😊😊	😊😊	😊	😊	😊	😊
Fork/Join	😊😊	😊😊😊😊	😊😊		😊😊😊 😊	😊😊😊😊

# Relationship between supporting program structure pattern and programming environment

	OpenMP	MPI	Java	Brook+/ CUDA	Cell
SPMD	😊😊😊	😊😊😊😊	😊😊	😊😊😊😊😊	😊😊😊😊
Loop Parallel	😊😊😊😊	😊	😊😊😊		😊😊😊
Master/ Slave	😊😊	😊😊😊	😊😊😊		😊😊😊😊
Fork/Joi n	😊😊😊		😊😊😊😊		😊😊

# Implementation Mechanism

